

COMS30030 - Image Processing and Computer Vision

Lecture 02

Filtering Images



Majid Mirmehdi | majid@cs.bris.ac.uk

Image Transforms

- **Image Transform** → a ***new representation*** of the input data (e.g. by re-encoding it in another [parameter] space, e.g. Fourier, Hough, Wavelet, Haar, etc.)
- Image Transforms are classic processing techniques, used in filtering, compression, feature extraction, and much more



Convolution

Convolution = A mathematical operation on two functions (f and g) that produces a third function ($h=f*g$), representing how the shape of one is modified by the other.

(Wikipedia)

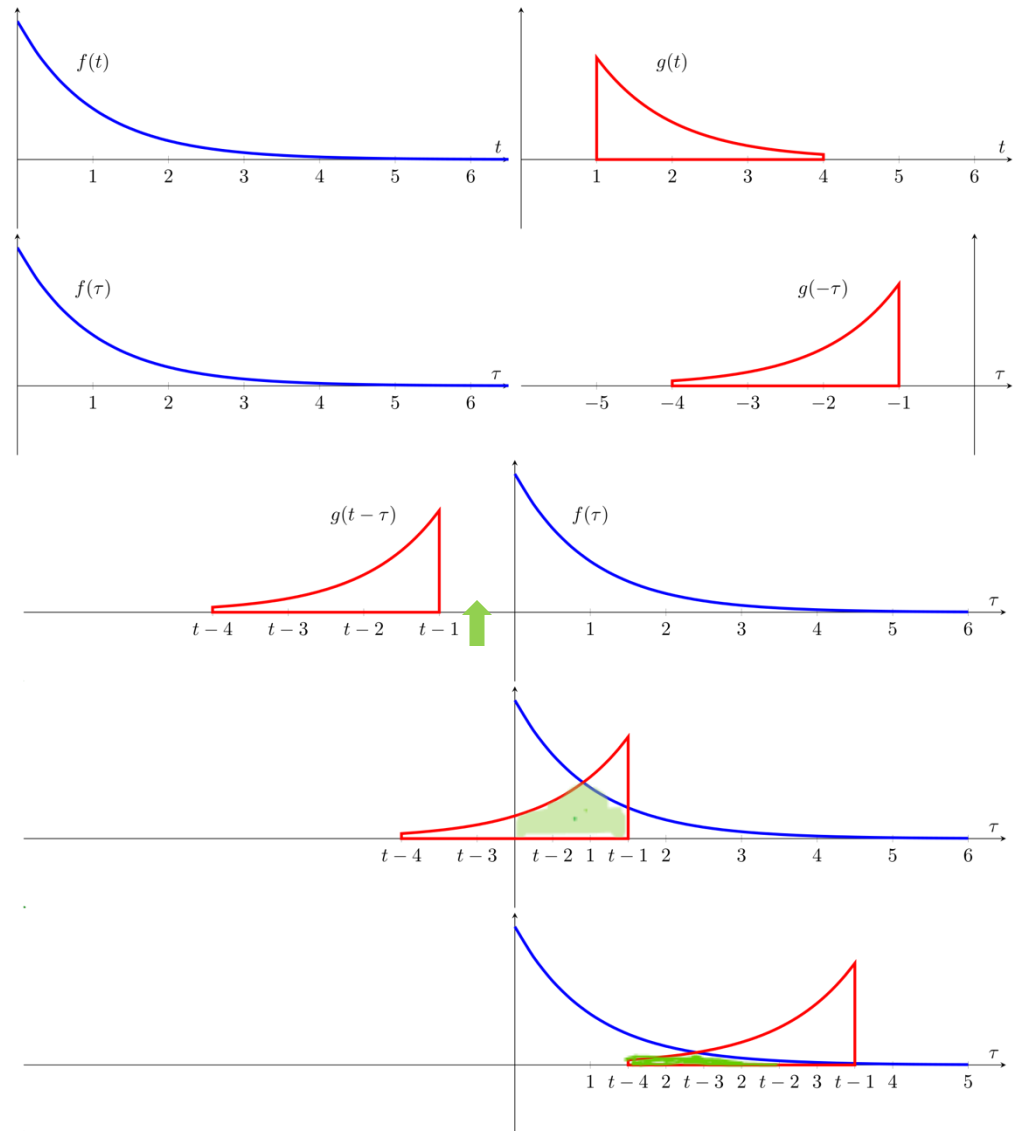
$$(f * g)(t) = \int_{-\infty}^{\infty} \underset{\text{Signal}}{f(t - x)} \overset{\text{Kernel}}{g(x)} \partial t$$

- Used for filtering images in the spatial domain
- Application in other parameter spaces, e.g. frequency domain
- Fundamental to Convolutional Neural Nets for deep Learning
- much, much more

An Intuitive look at Convolution


$$f * g = \int_{-\infty}^{\infty} f(t - \tau)g(\tau) \partial \tau$$

$$f * g = \int_{-\infty}^{\infty} f(\tau)g(t - \tau) \partial \tau$$




By Krishnavedala

An Intuitive look at Convolution



$$f * g = \int_{-\infty}^{\infty} f(t - \tau)g(\tau) \partial \tau$$



Autocorrelation

1D Discrete Convolution

$$g(x) = h(x) * f(x)$$

➤ f is the signal, h is the convolution filter

➤ h has an origin



Example 1D kernel

➤ Normalization factor (sum of the absolute values of the filter) is also part of the filter!

➤ The discrete version of convolution is defined

as:

$$g(x) = \sum_{m=-s}^s f(x-m)h(m) \quad \text{for } s \geq 1$$

2D Discrete Convolution

- The discrete version of 2D convolution is defined as:

$$g(x, y) = \sum_m \sum_n f(x - m, y - n)h(m, n)$$

		$y-1$	y	$y+1$	
$x-1$		43	12	61	
x		44	45	60	
$x+1$		43	50	61	

f *

	$n=-1$	$n=0$	$n=1$
$m=-1$	-1	0	1
$m=0$	-2	0	2
$m=1$	-1	0	1


h = **-68**

$$\begin{aligned}
 & f(x+1, y+1)h(-1, -1) \\
 & + f(x+1, y)h(-1, 0) \\
 & + f(x+1, y-1)h(-1, 1) \\
 & + f(x, y+1)h(0, -1) \\
 & + f(x, y)h(0, 0) \\
 & + f(x, y-1)h(0, 1) \\
 & + f(x-1, y+1)h(1, -1) \\
 & + f(x-1, y)h(1, 0) \\
 & + f(x-1, y-1)h(1, 1)
 \end{aligned}$$

2D Discrete Correlation

- The discrete version of 2D correlation is defined as:

$$g(x, y) = \sum_m \sum_n f(x + m, y + n) h(m, n)$$



Correlation \equiv Convolution
when kernel is symmetric
under 180° rotation, e.g.

1	2	1
1	2	1
1	2	1

Spatial Low/High Pass Filtering

Low Pass

$h =$

1	1	1
1	1	1
1	1	1



f




High Pass

$h =$

0	-1	0
-1	5	-1
0	-1	0



$f * h$



Properties of Convolution

- Commutative:

$$f * h = h * f$$

- Associative

$$(f * g) * h = f * (g * h)$$

- Distributes over addition

$$f * (g + h) = (f * g) + (f * h)$$

- Scalars factor out

$$af * g = f * ag = a(f * g)$$

- Identity:

Given unit impulse $e = [\dots, 0, 0, 1, 0, 0, \dots]$ then $f * e = f$

Image filtering example

$$h(m,n) = \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f(x,y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g(x,y)$

$$g(x,y) = \sum_m \sum_n f(x+m, y+n) h(m,n)$$

Credit: S. Seitz

Image filtering example

$$h(m,n) = \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f(x, y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g(x, y)$

	0	10							

$$g(x, y) = \sum_m \sum_n f(x + m, y + n) h(m, n)$$

Credit: S. Seitz

Image filtering example

$$h(m,n) \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f(x,y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g(x,y)$

	0	10	20						

$$g(x,y) = \sum_m \sum_n f(x+m, y+n) h(m,n)$$

Credit: S. Seitz

Image filtering example

$$h(m,n) = \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f(x,y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g(x,y)$

	0	10	20	30					

$$g(x,y) = \sum_m \sum_n f(x+m, y+n) h(m,n)$$

Credit: S. Seitz

Image filtering example

$$h(m, n) = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$f(x, y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g(x, y)$

	0	10	20	30	30				

$$g(x, y) = \sum_m \sum_n f(x + m, y + n) h(m, n)$$

Credit: S. Seitz

Image filtering example

$$h(m,n) \quad \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f(x,y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g(x,y)$

	0	10	20	30	30				

$$g(x,y) = \sum_m \sum_n f(x+m, y+n)h(m,n)$$

Credit: S. Seitz

Image filtering example

$$h(m,n) \quad \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f(x,y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g(x,y)$

	0	10	20	30	30				
						?			
				50					

$$g(x,y) = \sum_m \sum_n f(x+m, y+n) h(m,n)$$

Credit: S. Seitz

Image filtering example

$$h(m, n) = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$f(x, y)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g(x, y)$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$g(x, y) = \sum_m \sum_n f(x + m, y + n) h(m, n)$$

Credit: S. Seitz

Normalising the Convolution Result!

The weights will affect pixel values and the result could fall outside the 0-255 range → Normalise them such that they sum to 1.

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Smooth

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

3x3 Gaussian Blur

$$\frac{-1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & -476 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

Unsharp Masking

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$


Sharpen

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$



Basic Edge Detector

Gaussian Filter Kernel

The *Gaussian* filter is very commonly used in image processing. The parameter σ determines the width of the filter and hence the amount of smoothing.



$$g(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}$$




$$g(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$


Applying the Gaussian filter





Original




$\sigma = 1$



$\sigma = 5$



$\sigma = 10$




$\sigma = 30$

Source: Noah Snavely

Efficient Gaussian Filtering

Filtering with a 2D Gaussian can be achieved faster using two 1D Gaussian filters! This is because the linear Gaussian kernel is *separable*.



$$K = \frac{1}{256} \begin{pmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{pmatrix} = \frac{1}{16} \begin{pmatrix} 1 \\ 4 \\ 6 \\ 4 \\ 1 \end{pmatrix} \times \frac{1}{16} (1 \quad 4 \quad 6 \quad 4 \quad 1)$$

Filtering an $m \times m$ image with an $n \times n$ kernel is $\mathcal{O}(m^2 n^2)$ complexity,
but with a separable kernel, it would be $\mathcal{O}(m^2 n)$.
→ a significant reduction in computations!

Sources from fiveko.com and E.G.M. Petrakis

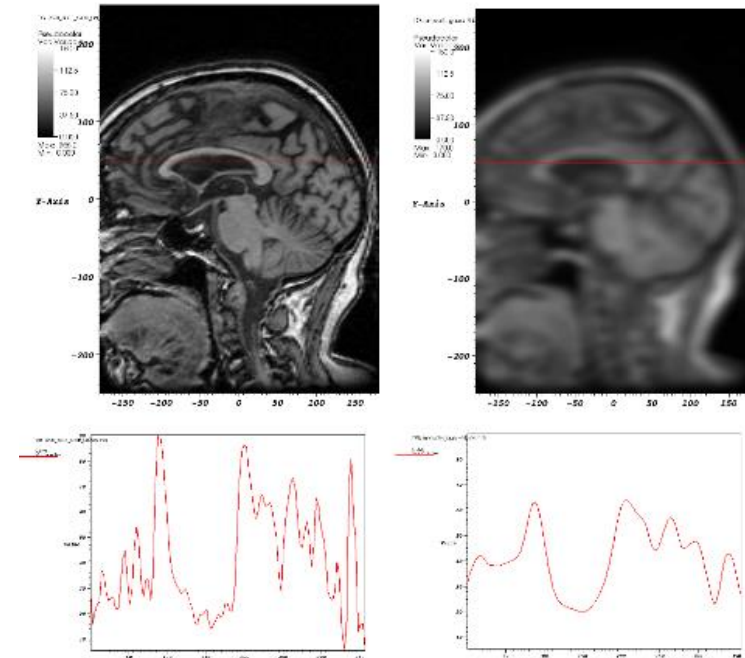
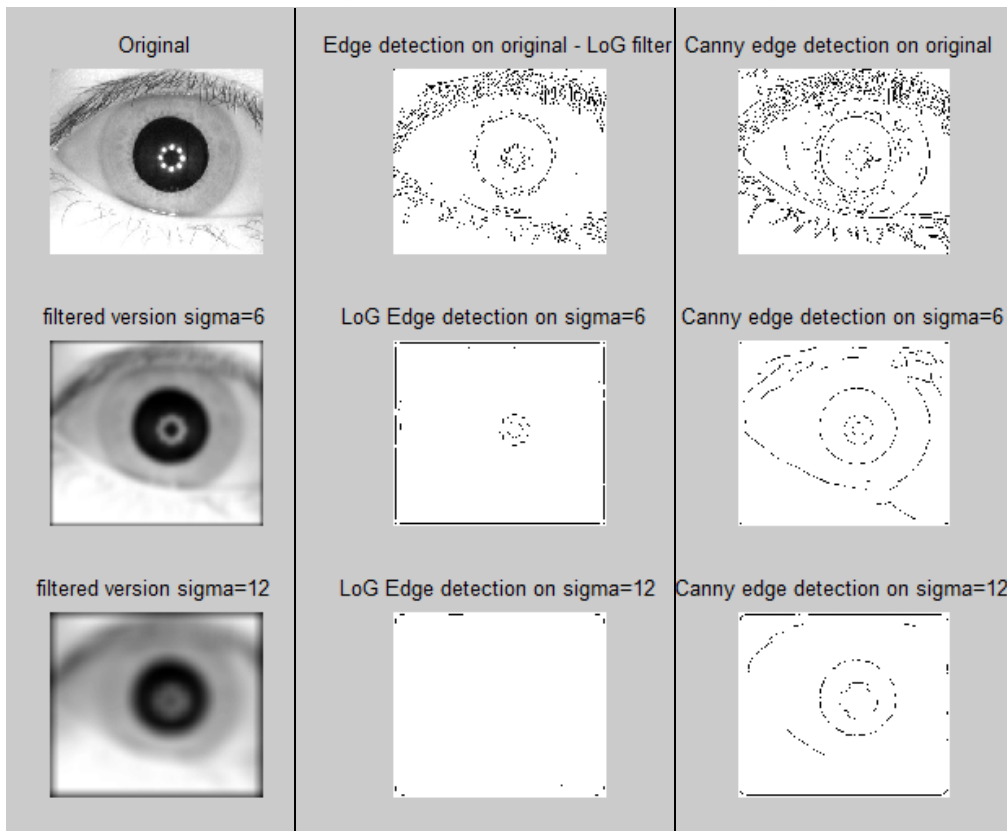
Separability of the Gaussian Filter

Separability means that a 2D convolution can be reduced to the product of two 1D convolutions - one on the rows and one on the columns:

$$\begin{aligned} g(x, y) &= \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{y^2}{2\sigma^2}} \right) \\ &= g(x) g(y) \end{aligned}$$

Example benefit of the Gaussian Filter

When we smooth an image, the easier the next stage of processing, in this case edge detection



Noise Removal: The Median Filter

- Returns the median value of the pixels in a neighborhood
- Relationship to a uniform blurring filter?
- Is non-linear

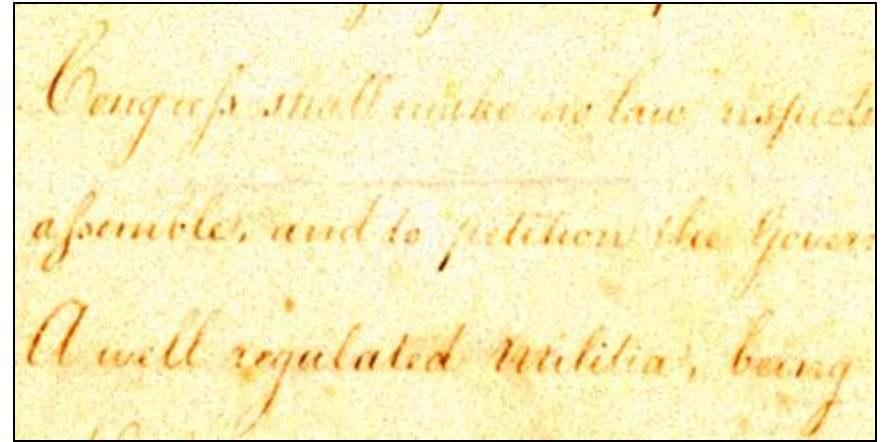
123	125	126	130	140
122	124	126	127	135
118	120	150	125	134
119	115	119	123	133
111	116	110	120	130

Neighbourhood values:

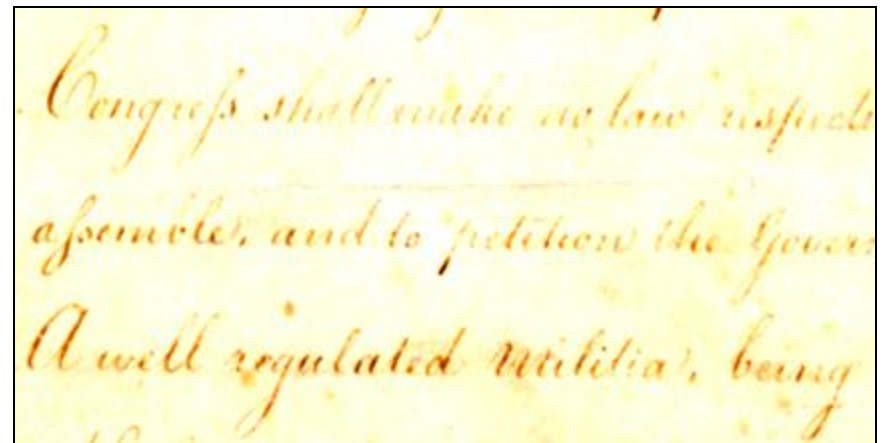
115, 119, 120, 123, 124,
125, 126, 127, 150

Median value: 124

original



median filtered



Slide adapted from Richard Peters

11

Noise Removal: The Median Filter


Original



salt & pepper
noise added



3x3
averaging
filter



3x3
median
filter



Noise Removal: The Median Filter

original



corrupted (*i.e.* $p=5\%$ that a bit is flipped)



median filtered




Slide adapted from Richard Peters

||


Median Filter: the algorithm

1. Let I be a monochrome image.
2. Let Z define a neighborhood of arbitrary shape.
3. At each pixel location, $\mathbf{p}=(x,y)$, in I ...
4. ... select the n pixels in the Z -neighborhood of \mathbf{p} ,
5. ... sort the n pixels in the neighborhood of \mathbf{p} by value into a list $L(j)$ for $j = 1, \dots, n$.
6. The output value at \mathbf{p} is $L(m)$, where $m = \lfloor n/2 \rfloor + 1$.

original



median filtered



Slide by Richard Alan Peters II

Sharpening Revisited

$h(m, n)$

0	-1	0
-1	4	-1
0	-1	0

Sharpening spatial filters seek to highlight fine detail

- Remove blurring from images
- Highlight edges

Sharpening filters are based on *spatial differentiation*.



Sharpening

What does blurring take away?



Let's add it back:




Photo credit: <https://www.flickr.com/photos/geezaweezer/16089096376/>


1st Derivative

The 1st derivative of a function is:

$$\frac{\partial f}{\partial x} = f(x+1) - f(x)$$

→ it's the difference between subsequent values and measures the rate of change of the function.


$f(x)$



5	5	4	3	2	1	0	0	0	6	0	0	0	0	1	3	1	0	0	0	0	7	7	7	7
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

	0	-1	-1	-1	-1	0	0	6	-6	0	0	0	0	1	2	-2	-1	0	0	0	7	0	0	0
--	---	----	----	----	----	---	---	---	----	---	---	---	---	---	---	----	----	---	---	---	---	---	---	---

$\frac{\partial f}{\partial x}$




2nd Derivative

The 2nd derivative of a function is:

$$\frac{\partial^2 f}{\partial x^2} = f(x+1) + f(x-1) - 2f(x)$$

→ Takes into account the values both before and after the current value.


$f(x)$



5	5	4	3	2	1	0	0	0	6	0	0	0	0	1	3	1	0	0	0	0	7	7	7	7
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

-1	0	0	0	0	1	0	6	-12	6	0	0	1	1	-4	1	1	0	0	7	-7	0	0
----	---	---	---	---	---	---	---	-----	---	---	---	---	---	----	---	---	---	---	---	----	---	---

$\frac{\partial^2 f}{\partial x^2}$



Using Second Derivatives For Image Enhancement

Now in 2-dimensions for images:

The 2nd derivative is more useful for image sharpening than the 1st derivative

- ✓ Stronger response to fine detail

The *Laplacian Filter*:

- ✓ Isotropic
- ✓ One of the simplest filters for sharpening

0	1	0
1	-4	1
0	1	0

The Laplacian

The Laplacian is defined as:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

where the partial 1st order derivative in the x direction is:

$$\frac{\partial^2 f}{\partial x^2} = f(x+1, y) + f(x-1, y) - 2f(x, y)$$

and in the y direction: $\frac{\partial^2 f}{\partial y^2} = f(x, y+1) + f(x, y-1) - 2f(x, y)$

So, the Laplacian can be:

$$\begin{aligned} \nabla^2 f = & [f(x+1, y) + f(x-1, y) \\ & + f(x, y+1) + f(x, y-1)] - 4f(x, y) \end{aligned}$$

0	1	0
1	-4	1
0	1	0

About the Laplacian Operator




As it is a derivative operator:

- it highlights gray-level discontinuities in an image.
- it de-emphasizes regions with slowly varying gray-levels.
- Very sensitive to noise as taking derivatives increases noise!

0	1	0	1	1	1
1	-4	1	1	-8	1
0	1	0	1	1	1
0	-1	0	-1	-1	-1
-1	4	-1	-1	8	-1
0	-1	0	-1	-1	-1

a
b c
d e


(a) Blurred image of the North Pole.
(b) Laplacian without scaling and (c) with scaling, (d) and (e) Image sharpening using two different masks.



LoG: Laplacian of Gaussian

Because the Laplacian is noise sensitive, it is always combined with a smoothing operation 😊

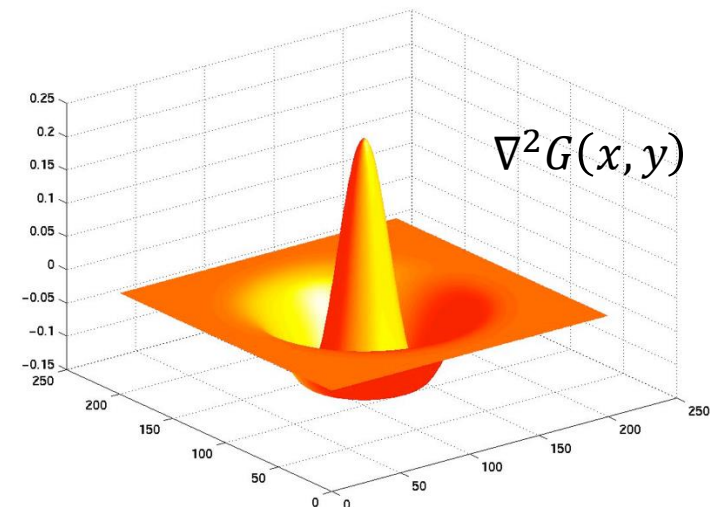
$$Output(x, y) = \nabla^2(f(x, y) * G(x, y))$$



$$\nabla^2(f(x, y) * G(x, y)) = \nabla^2 G(x, y) * f(x, y)$$

Laplacian of a
Gaussian-filtered
image

Laplacian of
Gaussian of a
filtered image



Summary: Image Filtering

- Images are often corrupted by random variations in intensity, illumination, or have poor contrast and cannot be used directly
- *Filtering* allows us to achieve:
 - *Enhancement*: improves contrast
 - *Smoothing*: removes noise
 - *Template matching*: detects known patterns
 - Feature Extraction: provides clues about objects etc., for further analysis
 - Many other uses...

Adapted from E.G.M. Petrakis